

# Deep Active Learning for Interactive 3D Segmentation Of Medical Images

---

Vincent Groff

September 10, 2018


# Table of contents

1. 3D Image Segmentation
2. Convolutional Neural Networks
3. Aims and Design

# 3D Image Segmentation

---


## 3D Medical Image Segmentation



**Figure 1:** Segmentation of the left atrium

**Main Problem:** Laborious, particularly in 3D

## 3D Medical Image Segmentation



**Figure 1:** Segmentation of the left atrium

**Main Problem:** Laborious, particularly in 3D

**How can we automate this?**

**Aim:** Separate the set of *background* pixels (**B**) from the set of *object* pixels (**O**)

**Define a cost function to be minimized over image X**

$$E(\mathbf{X}) = \sum_i R_i + \lambda \sum_i \sum_j B_{ij} \quad (1)$$

- $R_i$  - **Regional Penalty Term** - penalty for assigning pixel  $i$  to either the background (**B**) or object (**O**)
- $B_{ij}$  - **Boundary Penalty Term** - penalty depending on properties of pixels  $i$  and  $j$
- $\lambda > 0$  - a relative weighting

$R_i$  - **Regional Penalty Term**

Map  $(0,1)$  probabilities to  $(\infty, 0)$  penalties with negative log



## $R_i$ - Regional Penalty Term

Map (0,1) probabilities to  $(\infty, 0)$  penalties with negative log

$$R_i = \begin{cases} -\log P(i \in O) & i \in \mathbf{O} \\ -\log P(i \in B) & i \in \mathbf{B} \end{cases} \quad (2)$$

$B_{ij}$  - **Boundary Penalty Term**

Want to encourage boundaries at discontinuities in intensity

**Penalize boundaries that are placed between similar pixels**

# Conditional Random Fields

## $B_{ij}$ - Boundary Penalty Term

Want to encourage boundaries at discontinuities in intensity

**Penalize boundaries that are placed between similar pixels**

$$B_{ij} = \begin{cases} 0 & y_i \neq y_j \\ \frac{1}{r_{ij}} e^{-\frac{(x_i - x_j)^2}{2\sigma^2}} & y_i = y_j \end{cases} \quad (3)$$

- $X_i$  - grayscale value at pixel  $i$
- $y_i$  - binary label for pixel  $i$  (0 means  $i \in \mathbf{B}$ )
- $r_{ij}$  - euclidean distance between pixels  $i$  and  $j$

# Conditional Random Fields

## How do we solve it?

Represent the image as a graph with

- One node per pixel
- One source node (object)
- One sink (background)

# Conditional Random Fields

## How do we solve it?

Represent the image as a graph with

- One node per pixel
- One source node (object)
- One sink (background)

and

- Edges between pixels have a cost  $\lambda \frac{1}{r_{ij}} e^{\frac{-(x_i - x_j)^2}{2\sigma^2}}$
- Edges between pixels and the source terminal cost  $-\log P(i \in B)$
- Edges between pixels and the sink terminal cost  $-\log P(i \in O)$

**A minimum cost cut** which severs the source from the sink will now minimize the cost function

# Conditional Random Fields

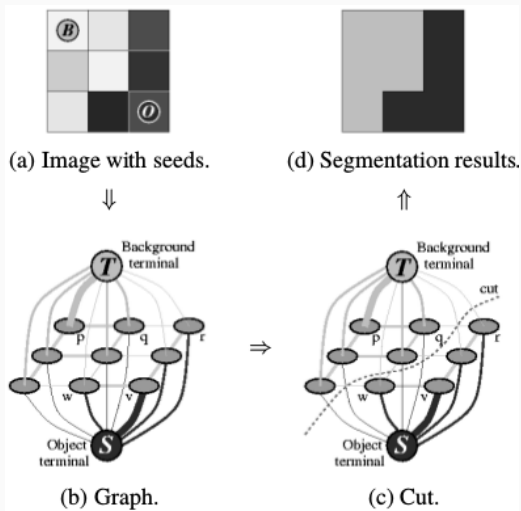


Figure 2: Minimum cost cut

**Maxflow:** Finding the maximal flow from source to sink

The edges that are saturated in maxflow are those in the minimal cost cut.

**Maxflow:** Finding the maximal flow from source to sink

The edges that are saturated in maxflow are those in the minimal cost cut.

Solve with **Breadth First Search**, by gradually saturating edges



**How do we get  $P(i \in B)$  and  $P(i \in O)$ ?**

**User interaction**

- User gives seed points
- Probability distribution produced from seed points
- Probabilities of other pixels inferred from distribution

**Works well, but...**

**How do we get  $P(i \in B)$  and  $P(i \in O)$ ?**

**User interaction**


- User gives seed points
- Probability distribution produced from seed points
- Probabilities of other pixels inferred from distribution

**Works well, but...**

**Requires object and background pixels to have distinct intensity values**

# Graph Cuts

Requires object and background pixels to have distinct intensity values



**Figure 3:** Grabcut algorithm in difficult cases

**Requires object and background pixels to have distinct intensity values**

This is the case for most traditional segmentation methods

# Convolutional Neural Networks

---

## Advantages

- Spatial operations
- Learning
- Performance

## Disadvantages

- Computationally expensive
- Poor generalisation to unseen types
- Need large datasets - expensive for medical imagery

**Classification CNNs** - gradually extract larger features from smaller ones

**Classification CNNs** - gradually extract larger features from smaller ones

**U-Net** - gradually re-combine those features to produce a segmentation map



# Segmentation CNNs




Figure 4: The U-Net architecture

## Improvements since original U-Net:

- Residual blocks
- Spatial Dropout
- Segmentation Layers

# Segmentation CNNs




Figure 5: The Isensee et al. architecture

**Segmentation Layers** encourage earlier layers to produce segmentations

## Aims and Design

---

## **Aims:**

Minimize the disadvantages of CNNs so that they can be used for 3D image segmentation

## **Disadvantages**


1. Computationally expensive
2. Poor generalisation to unseen types
3. Need large datasets - expensive for medical imagery

## **Bounding-box Image-specific Fine-Tuning Segmentations (BIFSeg)**

Combines CNN and CRF

Leverages user corrections

1. Draw bounding box around area of interest
2. CNN and CRF give initial guess
3. User provides correction scribbles
4. CNN fine-tunes using corrections and gives corrected result



**Figure 6:** The steps in the BIFSeg framework. In *blue* is the ground truth, in *cyan* the initial guess, in *yellow* pixels relabelled as background and in *red* is the final result

Two things going on:

- **Combining CNN and CRF**

CNN softmax output used as probability values for calculating  $R_i$



## Two things going on:

- **Combining CNN and CRF**

CNN softmax output used as probability values for calculating  $R_i$

- **Fine-Tuning CNN using CRF**

1. CNN, CRF and scribbles produce segmentation
2. CNN trained to match segmentation
3. Repeat steps 1 and 2  $n$  times

## Fine-Tuning CNN using CRF

CNN loss function is **weighted** on a per-pixel basis during fine-tuning, such that the weight for pixel  $i$  is

$$w_i = \begin{cases} 0 & i \in \mathbf{U} \\ w & i \in \mathbf{S} \\ 1 & \text{otherwise} \end{cases} \quad (4)$$

- $\mathbf{U}$  - set of pixels for which the current segmentation is uncertain
- $\mathbf{S}$  - set of user scribbles

## Fine-Tuning CNN using CRF

CNN loss function is **weighted** on a per-pixel basis during fine-tuning, such that the weight for pixel  $i$  is

$$w_i = \begin{cases} 0 & i \in \mathbf{U} \\ w & i \in \mathbf{S} \\ 1 & \text{otherwise} \end{cases} \quad (4)$$

- $\mathbf{U}$  - set of pixels for which the current segmentation is uncertain
- $\mathbf{S}$  - set of user scribbles

Pixels in  $\mathbf{U}$  are either **geodesically** near a scribble of opposite label or have  $P_i$  near 0.5

**BIFSeg works on unseen object types** - though it needs a CNN that has learned **generalisable features**

Performs better and with less user interaction on seen object types

**BIFSeg works on unseen object types** - though it needs a CNN that has learned **generalisable features**

Performs better and with less user interaction on seen object types

## **Transfer learning**

1. Start by using a CNN trained in generalised segmentation
2. After  $n$  images have been segmented, add them to a dataset
3. Load a new CNN with the generalised CNN weights and train it on the dataset
4. Repeat steps 2 and 3 until all images have been segmented

## Framework:

1. Produce a small number  $n$  of segmentations with some generalized CNN
2. Add these to a dataset
3. Train a new CNN through transfer learning on said dataset
4. Use the new